

Rishav Raj

✉ rairishav221@gmail.com | 📞 +91-9798600997 | 📍 Hyderabad, India
🌐 rishavraj.info | 🐙 github.com/rishavraj221 | 🔗 linkedin.com/in/rishav-raj

SUMMARY

I take problems end to end: system design, proof of concept, UI, backend, cloud setup. Recent work is mostly making a language model answer real questions from a company's own documents and show where the answer came from.

EDUCATION

B.Tech, Indian Institute of Technology Madras

2019 to 2023

EXPERIENCE

Evernorth Health Services (Cigna Group), Hyderabad

Mar 2025 to Present

Software Engineer, AI and Cloud [read more](#)

- **The app.** A chat app that answers product questions for sales agents, sourced from our Salesforce knowledge base. Built from scratch with the team. I did the system design, the UI and the backend services.
- **Ingestion.** A Playwright crawler signs in with its own role and walks seed articles to sub links at depth k. A parser per attachment type for xlsx, docx, pdf and pptx, then a small model normalises everything to markdown. Cron re crawls nightly, a content hash decides what is re parsed, and product experts curate it all from their own UI. [read more](#)
- **Retrieval.** Started on top k chunks from an Azure AI index. Moved the curated side to OKF and gave a thinking model tools, so it opens what it wants and queries spreadsheets as real tables rather than chunking them. Fixed a whole class of aggregate and cross sheet questions. [read more](#)
- **Model evaluation.** A framework that runs one setting for a window and scores it on user feedback, latency and cost, so a switch and a rollback are the same move. Now extending it so retrieval, ranking and where it runs are knobs too. [read more](#)
- **Production.** Okta, OpenShift, PostgreSQL, model calls through the internal AICoE gateway. Token bucket at the gateway, prompt caching on the stable prefix to hold latency and cost down, and separate dashboards for the app and for ingestion. [read more](#)

Feb 2024 to Feb 2025

Flexday AI Hyderabad

Software Developer, AI and Cloud [read more](#)

- **RBAC service.** Multi tenant, Node and Express on MongoDB, owning users, roles and machine identities. Clients sat on AWS, Azure and their own hardware, so cloud IAM could not cover them. [read more](#)
- **Hybrid checks.** Coarse permissions in the JWT validate locally, resource level checks hit a central endpoint. A per tenant policy version lands a role change on the next request rather than at expiry. [read more](#)
- **Tropicana chatbot.** Customer facing, GPT-4 and Elasticsearch, on AWS and the client's own servers. Cleared their security review, firewall rules and VPN setup first.

2023 to 2024

Fastail Hyderabad

Software Engineer [read more](#)

- **Social commerce feed.** Reels from local brands with ordering in the feed. Delivery area picks the candidates before engagement ranks them, using GeoJSON polygons on a 2dsphere index cached per location cell. [read more](#)
- **Payments and POS.** Razorpay webhooks as the source of truth rather than the browser callback, with ranked states so a late event cannot move an order backwards. Orders push into the restaurant's Petpooja POS keyed on our order id, so a retry cannot print twice. [read more](#)
- Node on Lambda, MongoDB, Cloudinary. Two engineers, I led the system design and built most of the frontend.

PROJECTS

Kureita, node based AI video studio

[kureita.com](#) [read more](#)

- **Runner.** Wired nodes resolve in topological order. FastAPI on Lambda self invokes for a fresh time budget, since generation outlasts the API Gateway timeout. [read more](#)
- **Webhooks and billing.** Provider callbacks are ED25519 verified and idempotent, so a retry cannot double charge. 15k lines of Python. [read more](#)

SKILLS

Search and AI: RAG pipelines end to end, crawling and parsing a corpus, chunking, hybrid and graph search, re ranking, evals and LLM as judge, guardrails, Azure AI Search, Weaviate, OpenAI, Claude, Gemini

Backend: Python, FastAPI, Node.js, Express, TypeScript, MongoDB, PostgreSQL, Redis, Elasticsearch

Infrastructure: AWS (Lambda, ECS Fargate, EC2, S3, API Gateway, SQS, ElastiCache, CloudFront, Secrets Manager), OpenShift, Terraform, Docker, SAM, GitHub Actions, Auth0, Okta

System design: rate limiting (token bucket, sliding window), caching and invalidation, idempotency, queues for load levelling, sharding and replication, consistency trade offs, multi tenant isolation, RED metrics

BIT OF PERSONAL LIFE

Training most days, gym and runs and cycling. Badminton and swimming. Out with run clubs, cycling groups and weekend treks, marathons when the calendar allows, Hyrox still pending. Road trips and long rides in between. Happy to have company for any of it.